

Attention

- Intro:-

- Helps LLMs understand & generate text by focusing on the relevant parts of input sequences.
- Attention allows the model to weigh the importance of different words in a sentence, enabling better content understanding.

- Embeddings Recap:-

- Embedding converts words into numerical vectors in a high-dim space.
- Similar words are mapped to close points in this space.
- "Apple" might map close to "Orange" in the context of fruits but close to "Phone" in a technology context.

- Attention Mechanism Basics:-

- Self-Attention:- Each word in a sentence is compared w/ every other word to compute a relevance score.

- Key Components:-

- Query (Q) :- Represents the word we're focusing on.
- Key (K) :- Represent all words ~~with~~ focusing in the sentence.
- Value (V) :- The actual content of the words that are weighted.

• Mathematical Formulation:-

- Dot Product:-

$$\text{Similarity} = Q \cdot K$$

if vector have unit length, dot product & cosine similarity are equivalent.

- Cosine Similarity:-

$$\frac{Q \cdot K}{\|Q\| \times \|K\|}$$

- Scaled Dot-Product

$$\frac{Q \cdot K}{\sqrt{d_K}}$$

• Used to prevent large dot products when vectors are high-dim.

• Softmax Function:-

- Converts raw similarity scores into probabilities.

$$\text{Softmax}(x_i) = \frac{e^{x_i}}{\sum_j e^{x_j}}$$

Ensures that scores sum to 1
are positive.

• Attention Score Calculation:-

Steps:-

- Compute similarity scores using dot product or scaled dot product.
- Apply the softmax function to these scores.
- Multiply the result by the value vectors to get the final weighted content vector.

• Multi-Head Attention:-

- Uses multiple sets of Q, K, V matrices to capture different aspects of word

relationship.

• Process:-

• Each "head" perform attention independently

• Results are concatenated & linearly transformed to produce the final output.

• Advantages:-

Allows the model to focus on different parts of the sentence simultaneously.

• Value Matrix (V):-

Role:- Transform the content vector into a form that is optimal for the next word prediction.

Interaction:-

→ Keys & Queries are used to calculate attention weights.

→ Values are transformed using these weights to get the final representation.

Summary of the Attention Process:

- Compute Similarities:- Between words using Q & K
- Softmax Normalization:- Convert similarities into Normalization.
- Weighted Sum:- Use weights to compute a weighted sum of the value vectors.
- Output:- The resulting vector is used for further processing (e.g., predicting the next word).
- Application in Transformers:-

Attention is applied multiple times within a Transformer to iteratively refine the understanding of content.

Transformer Architecture:-

- Embedding → Positional Encoding → Multi-Head Attention → Feedforward Network.
- Multi-Head attention is required across layers to build a deep understanding of text.